

# Machine Learning: Course Project

Olga Chernytska

June, 2018

## **Topic**

Predicting bus time to arrival at stops using GPS data only.

## **Topic Importance and Application**

Many European cities have electronic scoreboards on every bus stops that show minutes to arrival by transport routes. It makes public transport more convenient. But such scoreboards are not so widely used in Ukraine. As at now, the best prediction of time to arrival is provided by EasyWay on their website/application [1]. However, their prediction system is based on Google Maps statistics and sometimes on average speed on particular route interval. Google Maps may not count for number of stops on the route, stop time or traffic light schedule, while average speed approach is rather static and is too 'average'.

My aim is to build model using GPS data. The results of the model can be used for developing electronic scoreboards on the bus stops as well as by EasyWay to improve their time to arrival predictions.

## **Data**

The data used is collected by friend of mine, Maksym Hontar, using tool [2]. Data contains information on 8 distinct routes for Lviv city that pass Ukrainian Catholic University. Time period: since 2017-07-18 to 2018-05-28. The main fields are:

- datetime request time
- routecode, routeid, routename
- state 1 if bus is on route
- vehicleid, vehiclename
- x, y longitude and latitude of the vehicle at the time of request

Request is made every minute; response returned contains all the vehicle coordinates that are currently on route. Time of request is assumed to correspond to time when vehicle is in the location returned in response. However, it is not always true - data can be sent with a lag. Lags are random and can be of any length. Unfortunately, I cannot deal with it any way.

For analysis I selected route A53 (Id: 1723724) as it has information on movements of 36 distinct vehicles, which is about 3.1 millions of data points spread during 11 months.

There are two more data files:

- coordinates of stops for route 1723724 split by direction;
- coordinates of route points for route 1723724 split by direction.

## Methodology

I slightly modified method, explained in [3, p.356]. Authors propose to predict travel time from location  $c$  to stop  $j$  ( $T_{cj}$ ) using the following formula:

$$T_{cj} = T_{ij} - T_{ic}$$

where  $T_{ij}$  is predicted travel time between stops  $i$  and  $j$ , and  $T_{ic}$  - time that bus already passed from stop  $i$  to its location  $c$ .  $T_{ic}$  is determined by subtracting departure time at stop  $i$  from the time when bus is on location  $c$ .  $T_{ij}$  is determined using artificial neural networks. Originally, the following features are used:  $X_1$  - time of the day,  $X_2$  - code for station  $i$ ,  $X_3$  - code for station  $j$ ,  $X_4$  - travel time from 0 to  $i$ . The output  $Y$  is predicted bus travel time to reach stop  $j$  from stop  $i$  -  $T_{ij}$ .

My modifications are the following:

- Model predicts travel time only between nearest stops. To predict time to arrival to stop that is far away, I sum up predictions for the travel time between all the nearest stops, that are on the path.
- Instead of artificial neural network, I used linear regression as baseline model and gradient boosting as improved one.
- I used the following features: season (winter, spring, summer, autumn), weekday, hour, minute (split by 15-minute intervals) and route interval (characterized by starting and ending stops). Time features were extracted from the time stamp when bus is on the starting stop. All features were converted to dummy variables.

## Preprocessing

The data contains GPS coordinates only. Model inputs are season, weekday, hour, minute and route interval. Model output and  $Y$ 's for training - travel time between nearest stops. So some preprocessing steps has to performed.

- Step 0. When split by month, it can be seen that path for route 1723724 changes in October-December 2017, but than comes back to normal (see Figure 0). So these months are excluded from analysis.
- Step 1. Find all continues sequences 'time - coordinates' when bus is on the route. First, sort data frame by vehicle id and datetime. Find sequences of 1's ( $state = 1$  means that bus is on the route). Label first data point in each sequence as 'start', last data point as 'end' and all data points in between as 'on route'. Let's call this sequence as a drive. And let's give a unique id to each drive. Figure 1 shows how data looks like after this step.
- Step 2. Select only full drives and assign them to directions ('there' and 'back', or 0 and 1). Full drive means, that drive starts and ends near the starting and ending stops for particular route and direction. To do this, I found the starting and ending stops coordinates for the routes and directions (coordinates for 4 distinct stops in total).
- For every drive I checked whether data point labeled as 'start' is near the start stop for direction 0 or direction 1 ('near' means no further than 300 m away). I name the data point as 0/1, or -1 - if it is far from all start stops. The same - for the ending

stops. As a result, for each drive I gained labeling 0/1/-1 for each 'start' and 'end' data points. Figure 2 shows how data looks like after this step.

Step 3. At the previous step, I gained labeling 0/1/-1 for each 'start' and 'end' data points. Then I split drives by the directions. If 'start' and 'end' data points for the particular drive correspond and equal to 0, we know that this drive is full drive in the direction 0; and for direction 1 - respectively. There were a lot of partial drives, but after this step, all the dirty data was excluded from analysis.

Clean data (only full drives) was shifted to new data frames: one data frame contains drives in the direction 0 and another - in direction 1. Example of the data frame for each direction is shown in the Figure 3.

Step 4. Remove drives, that have big intervals between two nearest data points. Some drives that are recognized as full data is so sparse that it is impossible to understand how bus moved. These drives we should exclude to preserve data quality.

Step 5. Add artificial data. We are adding starting and ending stop coordinates (as observations) for those drives that start/end not near starting/ending stops. We assume that bus was in these locations 15/30/45 seconds earlier/later depending on how far the current drive start/end from true stops.

Data is sparse, if no errors we receive bus locations once in a minute. Now it is hard to understand what exact time bus was near the stop. We want to decrease this interval to 15 seconds, by adding 3 middle points between each pair of points for the drive.

Step 6. Remove duplicated locations at the drive ends. Sometimes state of the bus is 1 (which means that it is on route), but actually it is staying on the starting/ending stop with no change in location. For starting point we want to leave just the last one duplicate, because that represents time when the bus starts its move to next stop. For ending stop - opposite; we leave just the first duplicate as it is the time when bus ends its move from previous stop.

Step 7. Find closest coordinate on the drive to each stop. Using stops coordinates, for every direction and drive we can find at what exact date stamp bus was near the particular stop. We are looping through the every drive, looking for all data points that have minimal distances to all stops. But algorithmically, it is done a bit more efficiently. As a result, we have approximated time when the bus was near the every stop performing particular drive. Result are show in Figure 4.

Step 8. Rearrange data. In this step we rearranging data by calculating time in seconds between every two nearest stops for every drive. Results are shown in Figure 5.

Step 9. Test data. Firstly, we want to exclude drives for which algorithm couldn't find all the stops. This is important because time to move between stops are calculated incorrect in this case for all the intervals on the route. Secondly, route points that were recognized as closest to stops are highly spread around actual stop coordinates. We want to clean it as well by introducing a threshold - maximal distance by which coordinates for

observations can be far away from actual stop coordinates. Small threshold excludes more data points, big threshold leaves more dirty data. We selected threshold such way that most of the stops have non-overlapping surroundings and no more than quarter of observations is lost.

Now data looks better but still there are issues with couple of stops. Issues may occur because of dirty data on stop coordinates for the route. For instance, buses that moves in the direction 0, do not start on the 'true' starting point (highest one on the figure). This issue makes it impossible to predict time between start point and the next one.

The second problem that number of observations decreased unevenly for different route intervals. Now we have not so many observations for several intervals that may decrease the quality of prediction.

Figure 8 and Figure 9 show distribution of time to move between nearest stops. Data still has outliers, but most of the observations are located in narrow range close to median.

## Modeling

Two data sets - for directions 0 and 1 were concatenated by introducing new variable - 'direction'. This is done to train a single model, but not two models for every direction. Data set contains 150,864 data point.

Time variable contains important features, which were used to train a model - season (winter, spring, summer, autumn), weekday, hour and minute (by 15-minute intervals, four levels in total). Starting and ending stops were merged to be a single variable - 'route interval'. All the explanatory variables were converted to dummy variables; first level was excluded to prevent linear dependence of features. As a result, there are 57 explanatory variables. Variable to predict - travel time between every nearest stops, in seconds.

All the observations were randomly split into train and test set, 75% and 25% observations respectively. Metric to evaluate model was chosen to be median absolute error, because data set contains a lot of outliers, but we are interested more to fit non-outliers.

Baseline model is linear regression. Its median absolute error on test set is 30.86 seconds. The first improvement was done by excluding outliers from train set. Outliers are observations that lie beyond time travel intervals [mean - std; mean + std]. Intervals are calculated for every direction and route interval. As a result, 6.8% of train set observation identified as outliers and excluded from train set. Test set outliers were not excluded. This approach lead to decrease in median absolute error on train set to 29.28 seconds. Was decided to use train set without outliers as for linear regression it showed lower error rate.

Gradient Boosting Regressor (**sklearn** package) showed better result with tuned parameters. Parameters were tuned according to procedure recommended by DataRobot [4]:

- Pick 'n estimators' as large as computationally possible.
- Tune 'max depth', 'learning rate', 'min samples leaf', and 'max features' via grid search.
- Increase 'n estimators' even more and tune 'learning rate' again holding the other parameters fixed.

The actual results of Grid Search (`GridSearchCV` from `sklearn`) are the following:

- Number of estimators was fixed to 300. Grid Search was performed through parameters 'learning rate': [0.1, 0.05], 'max depth': [4, 6, 8], 'min samples leaf': [5, 9, 11]. The best combination - 'learning rate': 0.1, 'max depth': 6, 'min samples leaf': 11 - results in median absolute error of 19.93 seconds on cv set. It should be mentioned that cross validation was performed on train set with outliers excluded.

- As best 'min sample leaf' was selected to be 11, which is upper bound, I decided to additionally perform grid search through parameters: 'min samples leaf': [11, 14, 17, 20]. 'Learning rate' and 'max depth' were fixed at the level selected previously (0.1 and 6, respectively). Best 'min samples leaf' is still 11.

- Number of estimators was increased to 1000; 'max depth' = 6, 'min samples leaf' = 11. Search was performed through 'learning rate': [0.1, 0.5, 0.02, 0.01]. The best learning rate is 0.1. Median absolute error on cv set is 20.08 seconds.

Final model is Gradient Boosting Regressor with parameters: 'n estimators' = 1000, 'learning rate' = 0.1, 'max depth' = 6 and 'min samples leaf' = 11. Median absolute error on test set is 21.77 seconds, which is 40% decrease compared to baseline model.

Errors in various 'direction - route interval' vary a lot. Figure 10 and Figure 11 shows some groups with 25% higher and 25% lower group median error than total median error (21.77). The major issues are with 'direction = 0, route interval = #26 - #27', where median absolute error is 108.14 seconds, and with group 'direction = 1, route interval #27 - #28', which median error is 88.73 seconds.

## Demonstration

The task is to predict time to arrival to next 5 stops using time, direction and bus GPS coordinates [latitude, longitude] as inputs.

For demonstration purposes, I implemented function that randomly generates inputs - time, direction and GPS coordinates. Direction is a random integer - 0 or 1. Time is randomly selected from the data set available. GPS coordinates are generated from the route points data. I preprocessed route points data, so it is possible to get a point on every location on the route.

On the first step, algorithm identifies between what nearest stops bus is located. Location is represented by previous and next stops. Additionally, algorithm outputs share of interval between these stops that bus already passed (variable 'passed'). At most 5 next stops are returned as well.

On the second step, algorithm predicts time to arrival (in seconds) using model developed in previous section. Generated inputs are converted to appropriate format that model accepts - all the variable are converted to dummies with the same names as in train set. Travel time is calculated for every nearest pair of stops, time to arrival is calculated as cumulative sum for all the previous stops.

Predictions are done in the following way:

- To calculate travel time to first nearest stop, inputs are: previous and next stops, time and direction. The output is multiplied by (1- 'passed'), to subtract time that bus already passed from previous stop. Here travel time is equal to time to arrival to first nearest stop, which we predict.

- To calculate travel time from first nearest stop to second nearest stop (and for further route intervals), inputs are: first nearest and second nearest stops, time (but adding travel time calculated on previous step) and direction. I decided to update time because it can take 20 and more minutes when bus reaches last predicted 5th stop, but during this time rush hour may end or start, so travel time will be different. To calculate time to arrival at particular stops, algorithm cumulatively adds all the travel times to previous stops.

Algorithm outputs data frame that contains 5 next stops that bus is going to pass, seconds to arrival and actual arrival time for every stop.

Demonstration can be found on the Figure 13 and in Jupyter Notebook 'ML Project - Demonstration'.

### **Conclusion and Further work**

Throughout this project the following tasks were completed:

- Data was transformed from data frame 'time - vehicle id - latitude - longitude' to 'drive id - starting stop - ending stop - travel time in seconds'.

- Transformed data was used to train a model. Linear regression was selected to be a baseline model; it results in median absolute error of 30.86 seconds on test set. Baseline model was improved by 40% based on median absolute error decrease to 21.77 seconds on test set. This was done by excluding outliers from train set, changing model to be Gradient boosting regressor and tuning its parameters using Grid search.

- Developed working demo, that accepts bus GPS location, direction and time as inputs, and produces data frame that contains seconds to arrival and arrival time for the next 5 stops on the route.

The model was developed using only route id = 1223724. However, most of the code is standardized and can be easily converted to work for other routes. Demo works fast, because it loads and uses previously trained model, so it can be used in real time for future predictions. Median error of 21.77 seconds may be insignificant for customers - bus passengers.

However, accuracy of the model can be increased. Significant improvements lie in how data collected and processed. Now data is collected every minute, decreasing this interval to 15 seconds makes it possible not to use artificially generated midpoints. Approximation algorithm, that is expected to identify time when bus passes particular stop, works with errors which can be evaluated from Figure 6 and Figure 7. The other issue is mistakes in stops locations. Empirically, it was explored that buses moving in the direction 0, do not pass stop 0.

Another idea is to add variable that represent how the bus moved through previous route intervals during this drive. For instance, low speed on the previous route intervals may indicate that working day is a holiday.

As at now, model uses direction as an input. However, original data does not contain this variable. There are ways to deal with this problem: 1) use another source of data, for instance, EasyWay API; 2) develop algorithm that will identify direction; sequence of locations should be used as inputs in this case, because it is impossible to identify direction based on just single location.

## References:

1. EasyWay. <https://www.eway.in.ua/ua/cities/kyiv>
2. City Transport/Routes  
Links: [82.207.107.126:13541/SimpleRIDE/Lad/Sr/](https://82.207.107.126:13541/SimpleRIDE/Lad/Sr/),  
[82.207.107.126:13541/SimpleRIDE/LODA/Sr/](https://82.207.107.126:13541/SimpleRIDE/LODA/Sr/)
3. Wei Fan, Zegeye Gurmu. Dynamic Travel Time Prediction Models for Buses Using Only GPS Data  
Link: <https://www.sciencedirect.com/science/article/pii/S204604301630168X>
4. Mark Steadman. Gradient Boosted Regression Trees: DataRobot blog.  
Link: <https://blog.datarobot.com/gradient-boosted-regression-trees>

APPENDIX A. Figures: Preprocessing Stage.

Figure 0. Raw data for route 1723724, split by month.

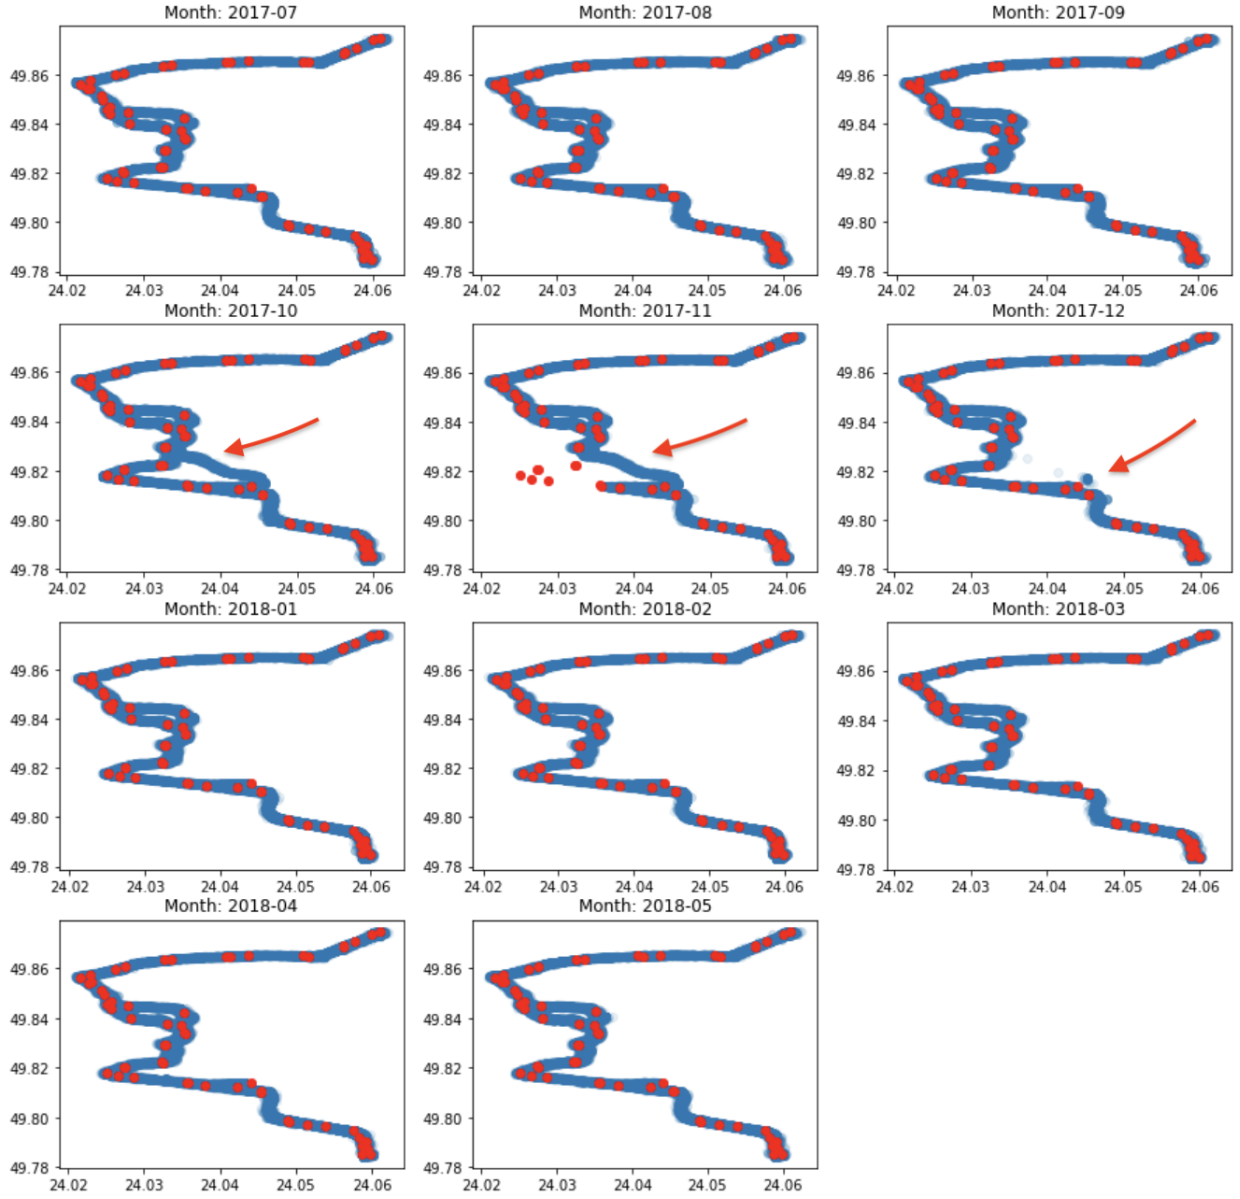




Figure 1. Results after Step 1.

| datetime                   | vehicle_id | state | lon       | lat       |               |
|----------------------------|------------|-------|-----------|-----------|---------------|
| 2017-07-20 05:13:21.811899 | 37984.0    | 0.0   | 24.058700 | 49.790617 | unknown       |
| 2017-07-20 05:14:31.32854  | 37984.0    | 0.0   | 24.059367 | 49.784700 | unknown       |
| 2017-07-20 05:15:33.007255 | 37984.0    | 0.0   | 24.059883 | 49.785050 | unknown DRIVE |
| 2017-07-20 05:16:22.452177 | 37984.0    | 1.0   | 24.059817 | 49.785450 | start         |
| 2017-07-20 05:17:21.934757 | 37984.0    | 1.0   | 24.059383 | 49.786133 | on route      |
| 2017-07-20 05:18:21.703389 | 37984.0    | 1.0   | 24.059183 | 49.787417 | on route      |
| 2017-07-20 05:19:31.559794 | 37984.0    | 1.0   | 24.059200 | 49.790567 | on route      |
| 2017-07-20 05:20:22.693668 | 37984.0    | 1.0   | 24.059067 | 49.791100 | on route      |
| 2017-07-20 05:21:31.948823 | 37984.0    | 1.0   | 24.058750 | 49.793283 | on route      |
| 2017-07-20 05:22:21.945059 | 37984.0    | 1.0   | 24.056267 | 49.795183 | on route      |
| 2017-07-20 05:23:22.38149  | 37984.0    | 1.0   | 24.053983 | 49.796300 | on route      |
| 2017-07-20 05:24:21.4714   | 37984.0    | 1.0   | 24.052783 | 49.796950 | on route      |
| 2017-07-20 05:25:22.271987 | 37984.0    | 1.0   | 24.048983 | 49.798917 | on route      |
| 2017-07-20 05:26:22.032683 | 37984.0    | 1.0   | 24.048883 | 49.798950 | on route      |
| 2017-07-20 05:27:21.796592 | 37984.0    | 1.0   | 24.046617 | 49.805717 | on route      |
| 2017-07-20 05:28:21.507506 | 37984.0    | 1.0   | 24.046833 | 49.808700 | end           |
| 2017-07-20 05:29:21.906135 | 37984.0    | 0.0   | 24.041800 | 49.811867 | unknown       |
| 2017-07-20 05:30:23.722561 | 37984.0    | 0.0   | 24.042267 | 49.812400 | unknown       |
| 2017-07-20 05:31:59.046659 | 37984.0    | 0.0   | 24.039600 | 49.813583 | unknown       |
| 2017-07-20 05:33:33.223789 | 37984.0    | 1.0   | 24.035483 | 49.813983 | start         |

Figure 2. Results after Step 2.

| datetime                   | vehicle_id | state | lon       | lat       |               |
|----------------------------|------------|-------|-----------|-----------|---------------|
| 2017-07-20 05:13:21.811899 | 37984.0    | 0.0   | 24.058700 | 49.790617 | unknown       |
| 2017-07-20 05:14:31.32854  | 37984.0    | 0.0   | 24.059367 | 49.784700 | unknown       |
| 2017-07-20 05:15:33.007255 | 37984.0    | 0.0   | 24.059883 | 49.785050 | unknown DRIVE |
| 2017-07-20 05:16:22.452177 | 37984.0    | 1.0   | 24.059817 | 49.785450 | start 1       |
| 2017-07-20 05:17:21.934757 | 37984.0    | 1.0   | 24.059383 | 49.786133 | on route      |
| 2017-07-20 05:18:21.703389 | 37984.0    | 1.0   | 24.059183 | 49.787417 | on route      |
| 2017-07-20 05:19:31.559794 | 37984.0    | 1.0   | 24.059200 | 49.790567 | on route      |
| 2017-07-20 05:20:22.693668 | 37984.0    | 1.0   | 24.059067 | 49.791100 | on route      |
| 2017-07-20 05:21:31.948823 | 37984.0    | 1.0   | 24.058750 | 49.793283 | on route      |
| 2017-07-20 05:22:21.945059 | 37984.0    | 1.0   | 24.056267 | 49.795183 | on route      |
| 2017-07-20 05:23:22.38149  | 37984.0    | 1.0   | 24.053983 | 49.796300 | on route      |
| 2017-07-20 05:24:21.4714   | 37984.0    | 1.0   | 24.052783 | 49.796950 | on route      |
| 2017-07-20 05:25:22.271987 | 37984.0    | 1.0   | 24.048983 | 49.798917 | on route      |
| 2017-07-20 05:26:22.032683 | 37984.0    | 1.0   | 24.048883 | 49.798950 | on route      |
| 2017-07-20 05:27:21.796592 | 37984.0    | 1.0   | 24.046617 | 49.805717 | on route      |
| 2017-07-20 05:28:21.507506 | 37984.0    | 1.0   | 24.046833 | 49.808700 | end 1         |
| 2017-07-20 05:29:21.906135 | 37984.0    | 0.0   | 24.041800 | 49.811867 | unknown       |
| 2017-07-20 05:30:23.722561 | 37984.0    | 0.0   | 24.042267 | 49.812400 | unknown       |
| 2017-07-20 05:31:59.046659 | 37984.0    | 0.0   | 24.039600 | 49.813583 | unknown       |
| 2017-07-20 05:33:33.223789 | 37984.0    | 1.0   | 24.035483 | 49.813983 | start         |

*Figure 3. Results after Step 3.*

| <b>drive_index</b> | <b>datetime</b>            | <b>lat</b> | <b>lon</b> |
|--------------------|----------------------------|------------|------------|
| 8                  | 2017-07-20 06:29:23.173075 | 49.872750  | 24.059667  |
| 8                  | 2017-07-20 06:30:24.701013 | 49.869133  | 24.056483  |
| 8                  | 2017-07-20 06:31:53.319531 | 49.865000  | 24.052017  |
| 8                  | 2017-07-20 06:33:40.596908 | 49.865100  | 24.051133  |
| 8                  | 2017-07-20 06:34:33.000431 | 49.865350  | 24.046867  |
| 8                  | 2017-07-20 06:35:33.301501 | 49.865050  | 24.042083  |
| 8                  | 2017-07-20 06:36:37.595833 | 49.864483  | 24.037167  |
| 8                  | 2017-07-20 06:37:29.130137 | 49.864150  | 24.035083  |
| 8                  | 2017-07-20 06:38:28.338238 | 49.863383  | 24.032600  |
| 8                  | 2017-07-20 06:39:32.822154 | 49.860300  | 24.027500  |
| 8                  | 2017-07-20 06:41:22.883643 | 49.856917  | 24.021283  |
| 8                  | 2017-07-20 06:42:27.378367 | 49.854450  | 24.022700  |
| 8                  | 2017-07-20 06:43:22.513019 | 49.854250  | 24.023033  |

Figure 4. Results after Step 7.

|   | drive_index |  | datetime                             | lat       | lon       | closest_stop       |
|---|-------------|--|--------------------------------------|-----------|-----------|--------------------|
| 0 | 8           |  | <u>2017-07-20 06:28:38.173075000</u> | 49.870820 | 24.057820 | stop #00: id=36865 |
| 1 | 8           |  | 2017-07-20 06:29:23.173075000        | 49.872750 | 24.059667 | none               |
| 2 | 8           |  | 2017-07-20 06:29:38.555059500        | 49.871846 | 24.058871 | none               |
| 3 | 8           |  | 2017-07-20 06:29:53.937044000        | 49.870942 | 24.058075 | stop #01: id=36853 |
| 4 | 8           |  | <u>2017-07-20 06:30:09.319028500</u> | 49.870037 | 24.057279 | none               |

Figure 5. Results after Step 8.

|   | drive_index | start_stop            | start_lat | start_lon | end_stop              | end_lat   | end_lon   | start_time                       | end_time                         | time_diff_seconds |
|---|-------------|-----------------------|-----------|-----------|-----------------------|-----------|-----------|----------------------------------|----------------------------------|-------------------|
| 0 | 8           | stop #00:<br>id=36865 | 49.870820 | 24.057820 | stop #01:<br>id=36853 | 49.870942 | 24.058075 | 2017-07-20<br>06:28:38.173075000 | 2017-07-20<br>06:29:53.937044000 | 75.763969         |
| 1 | 8           | stop #01:<br>id=36853 | 49.870942 | 24.058075 | stop #02:<br>id=37283 | 49.869133 | 24.056483 | 2017-07-20<br>06:29:53.937044000 | 2017-07-20<br>06:30:24.701013000 | 30.763969         |
| 2 | 8           | stop #02:<br>id=37283 | 49.869133 | 24.056483 | stop #03:<br>id=36822 | 49.865100 | 24.051133 | 2017-07-20<br>06:30:24.701013000 | 2017-07-20<br>06:33:40.596908000 | 195.895895        |
| 3 | 8           | stop #03:<br>id=36822 | 49.865100 | 24.051133 | stop #04:<br>id=36823 | 49.865125 | 24.043279 | 2017-07-20<br>06:33:40.596908000 | 2017-07-20<br>06:35:18.226233500 | 97.629326         |
| 4 | 8           | stop #04:<br>id=36823 | 49.865125 | 24.043279 | stop #05:<br>id=36821 | 49.864908 | 24.040854 | 2017-07-20<br>06:35:18.226233500 | 2017-07-20<br>06:35:49.375084000 | 31.148851         |

Figure 6. Step 9: Dirty and clean data for direction=0.




Figure 7. Step 9: Dirty and clean data for direction=1.




Figure 8. Distribution of time intervals: no zoom.


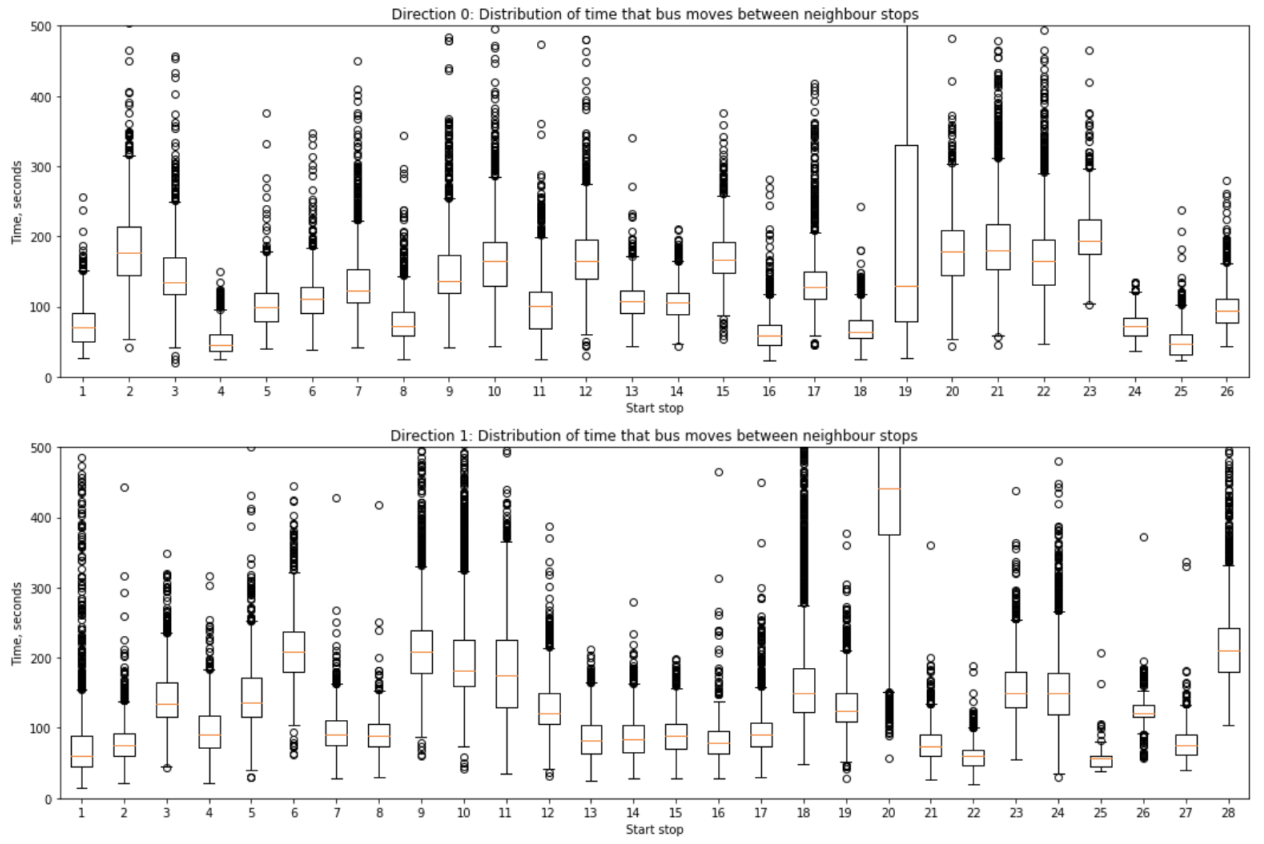


Figure 9. Distribution of time intervals: zoomed.



APPENDIX B. Figures: Modeling Stage.

Figure 10. Directions and route intervals with high median prediction error.

| direction | route_interval | n    | median_error | error_type |
|-----------|----------------|------|--------------|------------|
| 0         | #26_#27        | 807  | 108.142714   | high       |
| 1         | #27_#28        | 924  | 88.727332    | high       |
| 1         | #18_#19        | 932  | 43.556659    | high       |
| 0         | #02_#03        | 657  | 33.289124    | high       |
| 0         | #12_#13        | 745  | 31.618724    | high       |
| 0         | #10_#11        | 748  | 30.619026    | high       |
| 0         | #15_#16        | 614  | 30.306155    | high       |
| 1         | #13_#14        | 452  | 29.956692    | high       |
| 1         | #17_#18        | 940  | 29.701877    | high       |
| 1         | #00_#01        | 1035 | 28.627297    | high       |
| 1         | #04_#05        | 935  | 28.508142    | high       |
| 1         | #16_#17        | 912  | 28.413712    | high       |
| 1         | #25_#26        | 834  | 28.261078    | high       |
| 1         | #05_#06        | 858  | 28.031559    | high       |
| 0         | #11_#12        | 806  | 27.417206    | high       |



Figure 11. Directions and route intervals with low median prediction error.

| direction | route_interval | n   | median_error | error_type |
|-----------|----------------|-----|--------------|------------|
| 0         | #21_#22        | 418 | 16.309283    | low        |
| 1         | #01_#02        | 974 | 16.048940    | low        |
| 1         | #10_#11        | 634 | 15.712757    | low        |
| 1         | #23_#24        | 688 | 15.638464    | low        |
| 1         | #12_#13        | 417 | 15.081569    | low        |
| 0         | #19_#20        | 543 | 14.772386    | low        |
| 0         | #20_#21        | 571 | 13.429593    | low        |
| 0         | #25_#26        | 749 | 13.368453    | low        |
| 0         | #23_#24        | 657 | 13.253535    | low        |
| 0         | #04_#05        | 566 | 12.204376    | low        |
| 0         | #17_#18        | 349 | 12.052359    | low        |
| 1         | #11_#12        | 585 | 10.858193    | low        |
| 1         | #08_#09        | 50  | 8.884087     | low        |
| 1         | #07_#08        | 45  | 7.854191     | low        |

APPENDIX C. Figures: Demonstration.

Figure 12. Scheme of how algorithm works.

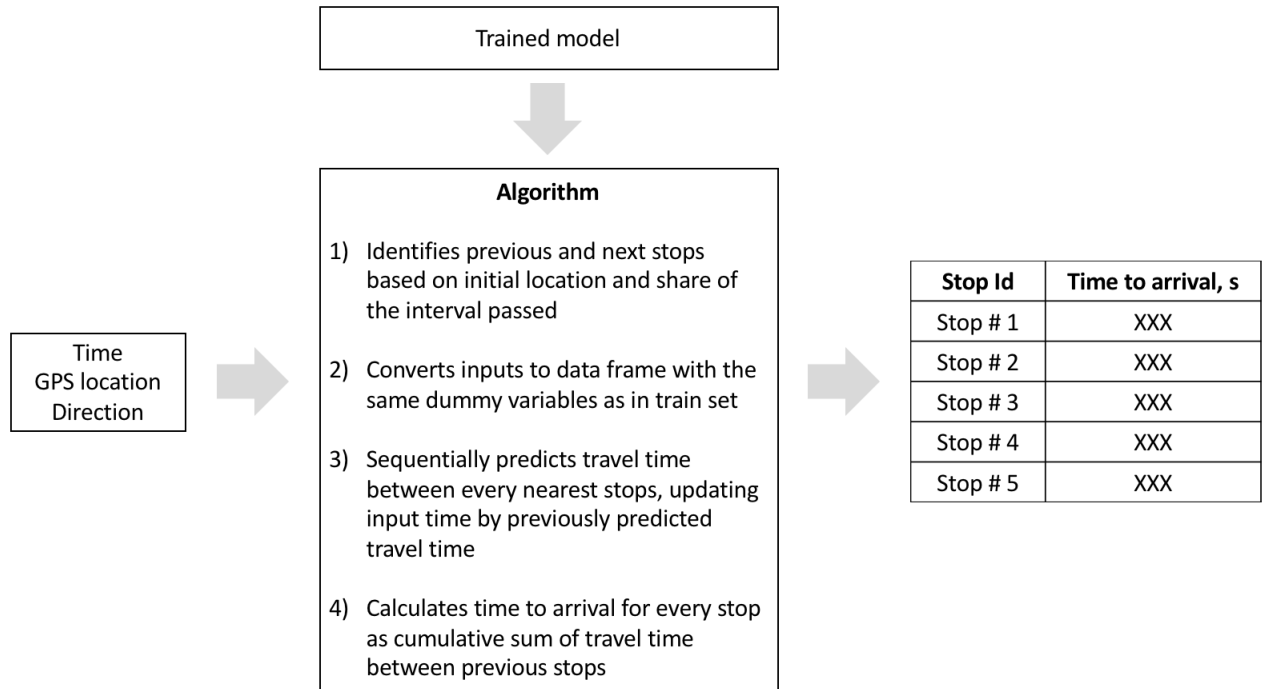



Figure 13. Demonstration.

Time: 2017-09-26T04:36:37.060401750  
 Direction: 1  
 Bus Location: [49.82356652125, 24.03342645875]

← randomly generated bus location,  
time and direction (input)



← visualization

previous and next stops for  
current bus location, share of  
interval passed

Bus is departed from stop #13: id=39164 and passed share=0.23 of the interval to stop #14: id=36559.

Next closest stops on the route: ['stop #14: id=36559' 'stop #15: id=36589' 'stop #16: id=36631'  
 'stop #17: id=36685' 'stop #18: id=36711']

|   | next_stop          | seconds_to_arrival | arrival_time                  |
|---|--------------------|--------------------|-------------------------------|
| 0 | stop #14: id=36559 | 155.847378         | 2017-09-26 04:39:12.907779750 |
| 1 | stop #15: id=36589 | 298.304957         | 2017-09-26 04:41:35.365358750 |
| 2 | stop #16: id=36631 | 433.488589         | 2017-09-26 04:43:50.548990750 |
| 3 | stop #17: id=36685 | 628.974608         | 2017-09-26 04:47:06.035010750 |
| 4 | stop #18: id=36711 | 795.359942         | 2017-09-26 04:49:52.420344750 |

← output

next 5 stops to be  
passed

APPENDIX C. Jupyter Notebooks attached.

1. ML Project Preprocessing.ipynb
2. ML Project Modeling.ipynb
3. ML Project Demonstration.ipynb